

04 - Salary Prediction

In this notebook, we will:

- Load the cleaned HR dataset.
- Prepare and encode features for regression.
- Split the data into training and testing sets.
- Train a simple **linear regression** model.
- We will then use **GridSearchCV** to test different tree based models, including **GradientBoosting**, **RandomForest** and **ElasticNet**.
- If more than 1 model shows promising results, we will also try to use a **stacking regressor** to combine models.
- We will try to find the best model to predict **MonthlyIncome**.
- Display sample predictions to illustrate how the model performs.

Since MonthlyIncome is a continuous variable, this task is approached as a regression problem.

Import Libraries & Load Data

We start by importing the necessary libraries and loading the cleaned dataset (saved as "hr_data_cleaned.csv").

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Regression libraries
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

# Set plotting style
sns.set(style="whitegrid")
plt.rcParams["figure.figsize"] = (12, 8)

# Load the cleaned dataset
df = pd.read_csv("hr_data_cleaned.csv")
print("DataFrame shape:", df.shape)
df.head()
```

DataFrame shape: (1470, 36)

	Age	Attrition	BusinessTravel	DailyRate	Department
0	41	Yes	Travel_Rarely	1102	Sales

1	49	No	Travel_Frequently	279	Research & Development
2	37	Yes	Travel_Rarely	1373	Research & Development
3	33	No	Travel_Frequently	1392	Research & Development
4	27	No	Travel_Rarely	591	Research & Development

	DistanceFromHome	Education	EducationField	EmployeeCount
EmployeeNumber \				
0	1	2	Life Sciences	1
1				
1	8	1	Life Sciences	1
2				
2	2	2	Other	1
4				
3	3	4	Life Sciences	1
5				
4	2	1	Medical	1
7				

	StandardHours	StockOptionLevel	TotalWorkingYears	\
0	80	0	8	
1	80	1	10	
2	80	0	7	
3	80	0	8	
4	80	1	6	

	TrainingTimesLastYear	WorkLifeBalance	YearsAtCompany	\
YearsInCurrentRole				
0	0	1	6	
4				
1	3	3	10	
7				
2	3	3	0	
0				
3	3	3	8	
7				
4	3	3	2	
2				

	YearsSinceLastPromotion	YearsWithCurrManager	TenureBucket
0	0	5	3-6
1	1	7	6-10
2	0	0	NaN
3	3	0	6-10
4	2	2	<3

```
[5 rows x 36 columns]
```

Simple Linear Regression with Preprocessing

Overview

We train a **Linear Regression model** to predict **MonthlyIncome**.

We use a pipeline

A **Pipeline** ensures smooth data preprocessing before training the model:

- **SimpleImputer**: Handles missing values in numerical features by filling them with the median. We actually do not have any missing values, hence we won't be using this.
- **OneHotEncoder**: Converts categorical features (university, degree, field of study) into numeric format for the model.

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler

# Target variable
target = "MonthlyIncome"

# Features
categorical_features = ["BusinessTravel", "Department",
                        "EducationField", "Gender", "OverTime", "JobRole", "MaritalStatus"]
numerical_features = ["JobLevel", "Age", "DistanceFromHome",
                      "Education", "NumCompaniesWorked", "TotalWorkingYears",
                      "TrainingTimesLastYear",
                      "YearsAtCompany", "YearsInCurrentRole",
                      "YearsSinceLastPromotion", "YearsWithCurrManager"]

# Drop rows where target is missing
df = df.dropna(subset=[target])

# Define preprocessing (One-Hot Encoding)
categorical_transformer = OneHotEncoder(handle_unknown="ignore")

full_transformer = ColumnTransformer(
    transformers=[
        ("num", StandardScaler(), numerical_features), # Scale
numerical features
        ("cat", categorical_transformer, categorical_features),
```

```

    ]
)

# Define Linear Regression model pipeline
linear_pipeline = Pipeline([
    ("preprocessor", full_transformer),
    ("model", LinearRegression())
])

# Train-test split
X = df[numerical_features + categorical_features]
y = df[target]
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Train the model
linear_pipeline.fit(X_train, y_train)

# Predictions
y_pred = linear_pipeline.predict(X_test)

# Evaluate model and print metrics

print(f"Linear Regression - Mean Squared Error:
{mean_squared_error(y_test, y_pred):.2f}")
print(f"Linear Regression - Mean Absolute Error:
{mean_absolute_error(y_test, y_pred):.2f}")
print(f"Linear Regression - R2 Score: {r2_score(y_test, y_pred):.2f}")

Linear Regression - Mean Squared Error: 1367382.92
Linear Regression - Mean Absolute Error: 886.77
Linear Regression - R2 Score: 0.94

```

GridSearchCV

- Exploring other models that possibly could be better than simple Linear Regression.
- Different models or better hyperparameters could improve performance.

Process

- Use `GridSearchCV` to test the following tree based models:
`DecisionTree`, `GradientBoosting`, `RandomForest`.
- Since we are only using tree based models, there is no need for a `StandardScaler`. Additionally, we have no missing values, eliminating the need for a `SimpleImputer`.
- Tune `max_depth`, `n_estimators`, and `learning_rate`.

- We find the best parameters for each model to get the most optimised result.
- Compare results and select the best-performing model.

```

from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import ElasticNet
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor,
GradientBoostingRegressor

# Target variable
target = "MonthlyIncome"

# Features
categorical_features = ["BusinessTravel", "Department",
"EducationField", "Gender", "OverTime", "JobRole", "MaritalStatus"]
numerical_features = ["JobLevel", "Age", "DistanceFromHome",
"Education", "NumCompaniesWorked", "TotalWorkingYears",
"TrainingTimesLastYear", "YearsAtCompany", "YearsInCurrentRole",
"YearsSinceLastPromotion", "YearsWithCurrManager"]

# Drop rows where target is missing
df = df.dropna(subset=[target])

# Define preprocessing (Handle missing values + One-Hot Encoding)
categorical_transformer = OneHotEncoder(handle_unknown="ignore")

full_transformer = ColumnTransformer(
    transformers=[
        ("cat", categorical_transformer, categorical_features),
    ]
)

# Define pipelines for different models
pipelines = {
    "DecisionTree": Pipeline([("preprocessor", full_transformer),
    ("model", DecisionTreeRegressor())]),
    "RandomForest": Pipeline([("preprocessor", full_transformer),
    ("model", RandomForestRegressor())]),
    "GradientBoosting": Pipeline([("preprocessor", full_transformer),
    ("model", GradientBoostingRegressor())])
}

# Define hyperparameter grids and find the best parameters
param_grids = {
    "DecisionTree": {
        "model__max_depth": [3, 5, 10, None], # Limits depth of tree
        (None = unlimited depth)
        "model__min_samples_split": [2, 5, 10], # Minimum samples
        required to split a node
    }
}

```

```

        "model__min_samples_leaf": [1, 2, 5], # Minimum samples
        required at a leaf node
        "model__max_features": [None, "sqrt", "log2"] # Number of
        features to consider when splitting
    },
    "RandomForest": {
        "model__n_estimators": [100, 200, 300],
        "model__max_depth": [None, 10, 20],
        "model__min_samples_split": [2, 5, 10]
    },
    "GradientBoosting": {
        "model__learning_rate": [0.01, 0.05, 0.1], # Lower values
        prevent overfitting
        "model__n_estimators": [100, 200, 300], # More trees improve
        performance
        "model__max_depth": [3, 5, 7], # Controls tree complexity
        "model__subsample": [0.7, 0.85, 1.0] # Subsample reduces
        variance
    }
}

# Train-test split (Ensure all years are represented)
X = df[numerical_features + categorical_features]
y = df[target]
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Fit and evaluate each model using GridSearchCV
results = {}
for name, pipeline in pipelines.items():
    print(f"Training {name}...")
    grid_search = GridSearchCV(pipeline, param_grids[name], cv=3,
scoring="neg_mean_squared_error", n_jobs=-1)
    grid_search.fit(X_train, y_train)

    # Get best model
    best_model = grid_search.best_estimator_
    y_pred = best_model.predict(X_test)

    # Evaluate
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    results[name] = {"MSE": mse, "R²": r2, "Best Params":
grid_search.best_params_}

# Print results
for name, result in results.items():
    print(f"{name} Regression - MSE: {result['MSE']:.2f}, R²:

```

```
{result['R²']:.2f}")
    print(f"Best Parameters: {result['Best Params']}\n")

Training DecisionTree...
Training RandomForest...
Training GradientBoosting...
DecisionTree Regression - MSE: 4754752.82, R²: 0.78
Best Parameters: {'model__max_depth': 5, 'model__max_features': None,
'model__min_samples_leaf': 5, 'model__min_samples_split': 2}

RandomForest Regression - MSE: 5146368.90, R²: 0.76
Best Parameters: {'model__max_depth': 10, 'model__min_samples_split':
10, 'model__n_estimators': 300}

GradientBoosting Regression - MSE: 4575549.94, R²: 0.79
Best Parameters: {'model__learning_rate': 0.05, 'model__max_depth': 3,
'model__n_estimators': 100, 'model__subsample': 0.7}
```

Analysis:

- All models seem to perform similarly well after multiple attempts at trial and testing the best possible parameters for each model.
- However simple linear regression seems to be a substantially better model.
- Perhaps we can try to use a stacking regressor to see if any model can assist LR?

Stacking Regressor

Now we try to :

- Combine Gradient Boosting, Random Forest, Decision Tree and Linear Regression using a Stacking Regressor.
- Use different stacking configurations to see which model combination performs best.

```
from sklearn.ensemble import StackingRegressor

# Base models. For gb, rf and DT, we use the optimised parameters.
rf_model = RandomForestRegressor(n_estimators=200, max_depth=10,
min_samples_split=10, random_state=42)
gb_model = GradientBoostingRegressor(n_estimators=100,
learning_rate=0.05, max_depth=3, subsample=0.7, random_state=42)
dt_model = DecisionTreeRegressor(max_depth=5, min_samples_leaf=5,
min_samples_split=2, random_state=42)
lin_model = LinearRegression()

# Stacking configurations and must include LR
stacking_models = {
    "Stacking_RF_GB_DT": StackingRegressor(
        estimators=[("rf", rf_model), ("gb", gb_model), ("dt",
dt_model)], final_estimator=LinearRegression())
```

```

    ),
    "Stacking_RF_DT_LR": StackingRegressor(
        estimators=[("rf", rf_model), ("dt", dt_model), ("lr",
lin_model)],
final_estimator=GradientBoostingRegressor(n_estimators=100,
learning_rate=0.05, random_state=42)
    ),
    "Stacking_GB_DT_LR": StackingRegressor(
        estimators=[("gb", gb_model), ("dt", dt_model), ("lr",
lin_model)], final_estimator=RandomForestRegressor(n_estimators=100,
random_state=42)
    ),
    "Stacking_RF_GB_DT_LR": StackingRegressor(
        estimators=[("rf", rf_model), ("gb", gb_model), ("dt",
dt_model), ("lr", lin_model)],
final_estimator=GradientBoostingRegressor(n_estimators=100,
learning_rate=0.05, random_state=42)
    ),
    "Stacking_RF_LR": StackingRegressor(
        estimators=[("rf", rf_model), ("lr", lin_model)],
        final_estimator=GradientBoostingRegressor(n_estimators=100,
learning_rate=0.05, random_state=42)
    ),
    "Stacking_GB_LR": StackingRegressor(
        estimators=[("gb", gb_model), ("lr", lin_model)],
        final_estimator=RandomForestRegressor(n_estimators=100,
random_state=42)
    ),
    "Stacking_DT_LR": StackingRegressor(
        estimators=[("dt", dt_model), ("lr", lin_model)],
        final_estimator=GradientBoostingRegressor(n_estimators=100,
learning_rate=0.05, random_state=42)
    )
}

```

Train and evaluate each stacking model

```

results = {}
for name, model in stacking_models.items():
    print(f"Training {name}...")
    pipeline = Pipeline([
        ("preprocessor", full_transformer),
        ("model", model)
    ])

    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

```



```

    results[name] = {"MSE": mse, "R²": r2}

# Print results
for name, result in results.items():
    print(f"{name} - MSE: {result['MSE']:.2f}, R²: {result['R²']:.2f}\n")

Training Stacking_RF_GB_DT...
Training Stacking_RF_DT_LR...
Training Stacking_GB_DT_LR...
Training Stacking_RF_GB_DT_LR...
Training Stacking_RF_LR...
Training Stacking_GB_LR...
Training Stacking_DT_LR...
Stacking_RF_GB_DT - MSE: 4528687.10, R²: 0.79

Stacking_RF_DT_LR - MSE: 4703000.71, R²: 0.78

Stacking_GB_DT_LR - MSE: 5485862.61, R²: 0.75

Stacking_RF_GB_DT_LR - MSE: 4751364.50, R²: 0.78

Stacking_RF_LR - MSE: 4675307.36, R²: 0.79

Stacking_GB_LR - MSE: 5392358.96, R²: 0.75

Stacking_DT_LR - MSE: 4606263.94, R²: 0.79

```

- Linear Regression seems to remain the best model based on its R^2 value. This was despite multiple trial and errors, including letting LR be the final estimator. This resulted in all models having an R^2 value of approx 0.8. In any case, standalone LR has an R^2 score of 0.94. Hence, we decided to just simply go ahead with that model for our predictions.

```

# Create a DataFrame with actual and predicted values
results_df = X_test.copy() # Copy test features for reference
results_df['actual_MonthlyIncome'] = y_test.values # Add actual target values
results_df['predicted_MonthlyIncome'] = y_pred # Add predictions

# View or save the result
results_df.head(20)

```

	JobLevel	Age	DistanceFromHome	Education	
NumCompaniesWorked \					
1041	2	28	5	3	0
184	2	53	13	2	1
1222	1	24	22	1	1

67	3	45	7	3	2
220	2	36	5	2	8
494	1	34	14	3	1
430	1	35	22	3	0
240	1	39	1	4	7
218	3	45	6	3	6
49	1	35	8	1	1
665	1	47	2	4	1
926	3	43	4	4	3
617	2	44	4	3	9
361	1	40	10	4	3
1423	1	22	1	2	0
1244	1	30	2	4	0
1250	2	29	1	3	8
752	1	36	16	4	1
271	3	47	29	4	1
1055	4	34	15	3	7

	TotalWorkingYears	TrainingTimesLastYear	YearsAtCompany	\
1041	6	4	5	
184	5	3	4	
1222	1	2	1	
67	25	2	1	
220	16	3	13	
494	8	3	8	
430	6	2	5	
240	7	1	3	
218	23	2	19	
49	1	2	1	
665	3	3	3	
926	23	3	21	
617	10	2	5	
361	10	3	7	
1423	4	2	3	

1244	10	2	9
1250	10	5	3
752	18	1	17
271	10	2	10
1055	16	3	14

YearsInCurrentRole	YearsSinceLastPromotion
YearsWithCurrManager \	
1041	4 1
3	
184	2 1
3	
1222	0 0
0	
67	0 0
0	
220	11 3
7	
494	2 0
6	
430	4 4
3	
240	2 1
2	
218	7 12
8	
49	0 0
1	
665	2 1
2	
926	7 15
17	
617	2 2
3	
361	7 1
7	
1423	2 1
2	
1244	7 0
7	
1250	2 0
2	
752	13 15
14	
271	7 9
9	
1055	8 6
9	

Gender \	BusinessTravel	Department	EducationField
1041 Male	Travel_Rarely	Sales	Medical
184 Female	Travel_Rarely	Research & Development	Medical
1222 Male	Travel_Rarely	Human Resources	Human Resources
67 Male	Travel_Rarely	Research & Development	Life Sciences
220 Male	Travel_Rarely	Research & Development	Life Sciences
494 Female	Travel_Rarely	Sales	Technical Degree
430 Male	Travel_Rarely	Research & Development	Life Sciences
240 Female	Travel_Rarely	Research & Development	Medical
218 Female	Non-Travel	Sales	Medical
49 Male	Travel_Rarely	Research & Development	Life Sciences
665 Female	Travel_Rarely	Sales	Life Sciences
926 Female	Travel_Rarely	Sales	Marketing
617 Male	Travel_Rarely	Research & Development	Medical
361 Female	Travel_Rarely	Research & Development	Life Sciences
1423 Male	Travel_Rarely	Research & Development	Life Sciences
1244 Female	Travel_Frequently	Research & Development	Technical Degree
1250 Male	Travel_Frequently	Research & Development	Life Sciences
752 Female	Travel_Rarely	Research & Development	Life Sciences
271 Male	Non-Travel	Research & Development	Life Sciences
1055 Male	Travel_Frequently	Research & Development	Medical

OverTime	JobRole	MaritalStatus
actual_MonthlyIncome \		
1041 No	Sales Executive	Single
8463		
184 No	Manufacturing Director	Divorced
4450		

1222	No	Human Resources	Married
1555			
67	No	Research Scientist	Divorced
9724			
220	No	Laboratory Technician	Single
5914			
494	Yes	Sales Representative	Divorced
2579			
430	No	Laboratory Technician	Single
4230			
240	No	Laboratory Technician	Divorced
2232			
218	No	Sales Executive	Single
8865			
49	No	Laboratory Technician	Married
2269			
665	Yes	Sales Representative	Single
3294			
926	No	Sales Executive	Single
10231			
617	No	Healthcare Representative	Single
5933			
361	Yes	Laboratory Technician	Married
2213			
1423	No	Research Scientist	Single
3375			
1244	No	Research Scientist	Single
4968			
1250	Yes	Healthcare Representative	Single
6294			
752	No	Laboratory Technician	Single
2743			
271	Yes	Manager	Married
11849			
1055	No	Research Director	Divorced
17007			

	predicted_MonthlyIncome
1041	7566.721853
184	7127.297313
1222	3927.388603
67	3118.635256
220	3118.635256
494	3118.635256
430	3118.635256
240	3118.635256
218	6444.223080
49	3118.635256
665	3132.513215

926	6491.975111
617	6240.055262
361	3118.635256
1423	3118.635256
1244	2971.171862
1250	6863.662739
752	3118.635256
271	16365.336705
1055	15067.341724

Feature Encoding & Train-Test Split

We also predicted **MonthlyIncome** using another set of predictors, for example, Age, DistanceFromHome, and OverTime. Then, we apply one-hot encoding to the categorical variables.

```
# Define our target variable and features
target = "MonthlyIncome"
features = ["Age", "DistanceFromHome", "OverTime", "BusinessTravel",
"JobLevel"]

# One-hot encode the categorical columns among the features
df_encoded = pd.get_dummies(df[features], drop_first=True)

# Prepare feature matrix X and target vector y
X = df_encoded.values
y = df[target].values

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

print("Training set shape:", X_train.shape)
print("Test set shape:", X_test.shape)

Training set shape: (1176, 6)
Test set shape: (294, 6)
```

Train a Random Forest Regressor

We also trained a Random Forest Regressor on our training data. This model is robust, handles non-linearities well, and can provide feature importance insights.

```
# Initialize and train the Random Forest Regressor
rf_regressor = RandomForestRegressor(n_estimators=100,
random_state=42)
rf_regressor.fit(X_train, y_train)

RandomForestRegressor(random_state=42)
```

Evaluate the Regression Model

We now predict MonthlyIncome on the test set and evaluate our model using metrics such as Mean Absolute Error (MAE), Mean Squared Error (MSE), and R² score.

```
# Predict MonthlyIncome on the test set
y_pred = rf_regressor.predict(X_test)

# Calculate evaluation metrics
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Absolute Error: {mae:.2f}")
print(f"Mean Squared Error: {mse:.2f}")
print(f"R^2 Score: {r2:.3f}")
```

```
Mean Absolute Error: 1100.05
Mean Squared Error: 2226383.00
R^2 Score: 0.898
```

Sample Predictions

Here, we select a few samples from the test set and compare the predicted MonthlyIncome against the actual values.

```
# Select 5 random samples from the test set
sample_indices = np.random.choice(X_test.shape[0], size=5,
replace=False)
sample_features = X_test[sample_indices]
sample_true = y_test[sample_indices]

# Predict using the model for these samples
sample_preds = rf_regressor.predict(sample_features)

# Create a DataFrame to display sample predictions
sample_df = pd.DataFrame(sample_features, columns=df_encoded.columns)
sample_df["TrueMonthlyIncome"] = sample_true
sample_df["PredictedMonthlyIncome"] = sample_preds
sample_df
```

	Age	DistanceFromHome	JobLevel	OverTime_Yes
0	28	8	1	False
1	28	1	1	True
2	44	4	5	False

3	58	21	4	True
False				
4	34	9	3	False
False				

	BusinessTravel_Travel_Rarely	TrueMonthlyIncome
PredictedMonthlyIncome		
0	True	3310
4215.476167		
1	False	2154
2821.833381		
2	True	19190
18899.860000		
3	True	17875
14858.790000		
4	True	8500
9434.635000		